



Assignment 5

Part 1: Decision implementation on odinal data

```
In [2]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OrdinalEncoder
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import classification_report, accuracy_score
import matplotlib.pyplot as plt

# Load car.data
cols = ["buying", "maint", "doors", "persons", "lug_boot", "safety", "class"]
df = pd.read_csv("car.data", names=cols)

print("Dataset shape:", df.shape)
df.head()
```

Dataset shape: (1728, 7)

```
Out[2]:   buying  maint  doors  persons  lug_boot  safety  class
0    vhigh    vhigh      2        2    small     low  unacc
1    vhigh    vhigh      2        2    small    med  unacc
2    vhigh    vhigh      2        2    small   high  unacc
3    vhigh    vhigh      2        2     med     low  unacc
4    vhigh    vhigh      2        2     med    med  unacc
```

```
In [3]: # Explicit ordinal ordering for each attribute
ord_maps = [
    ["vhigh", "high", "med", "low"],    # buying
    ["vhigh", "high", "med", "low"],    # maint
    ["2", "3", "4", "5more"],          # doors
    ["2", "4", "more"],                # persons
    ["small", "med", "big"],           # lug_boot
    ["low", "med", "high"]             # safety
]

X = df.iloc[:, :-1]
y = df.iloc[:, -1]

enc = OrdinalEncoder(categories=ord_maps)
X_encoded = enc.fit_transform(X)

print("Encoded feature sample:\n", X_encoded[:5])
```

```
Encoded feature sample:
```

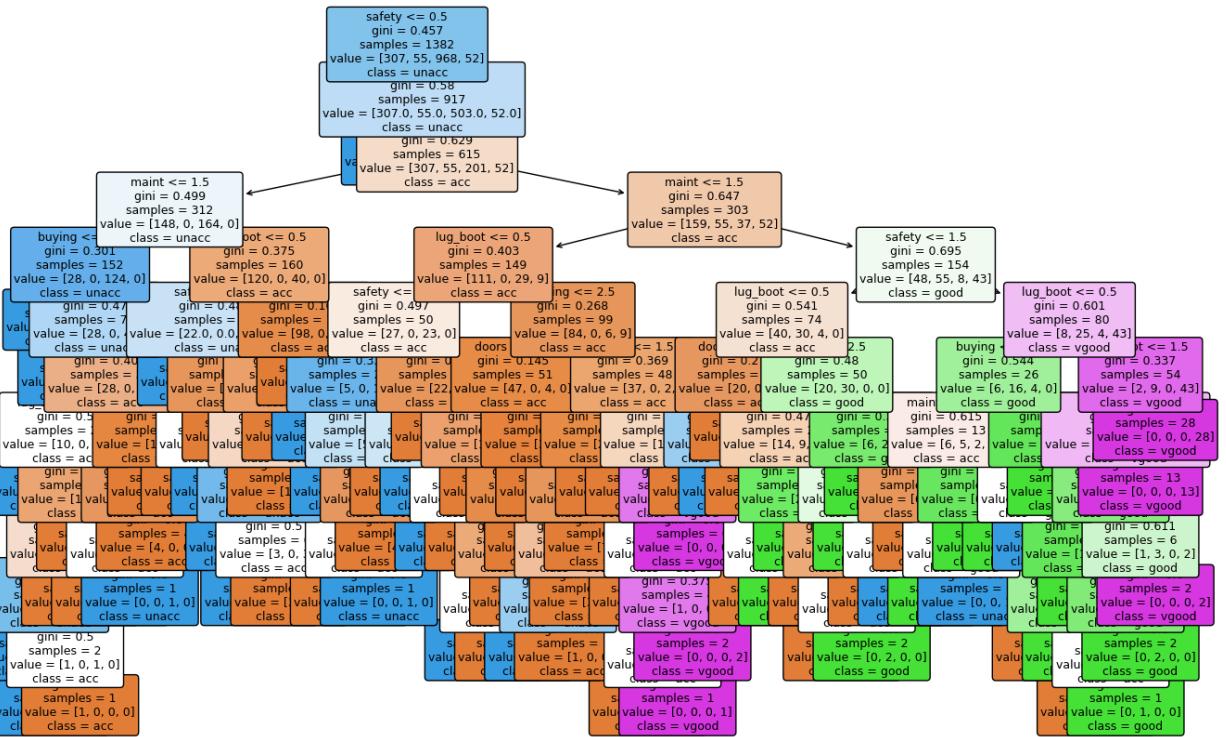
```
[[0. 0. 0. 0. 0. 0.]  
 [0. 0. 0. 0. 0. 1.]  
 [0. 0. 0. 0. 0. 2.]  
 [0. 0. 0. 0. 1. 0.]  
 [0. 0. 0. 0. 1. 1.]]
```

```
In [4]: X_train, X_test, y_train, y_test = train_test_split(  
         X_encoded, y, test_size=0.2, stratify=y, random_state=42  
     )  
  
# Fully grown decision tree  
clf = DecisionTreeClassifier(random_state=42)  
clf.fit(X_train, y_train)  
  
y_pred = clf.predict(X_test)  
  
print("Accuracy:", accuracy_score(y_test, y_pred))  
print(classification_report(y_test, y_pred))
```

```
Accuracy: 0.9913294797687862
```

	precision	recall	f1-score	support
acc	0.97	0.99	0.98	77
good	0.93	1.00	0.97	14
unacc	1.00	1.00	1.00	242
vgood	1.00	0.92	0.96	13
accuracy			0.99	346
macro avg	0.98	0.98	0.98	346
weighted avg	0.99	0.99	0.99	346

```
In [5]: plt.figure(figsize=(16, 10))  
plot_tree(  
    clf,  
    feature_names=cols[:-1],  
    class_names=clf.classes_,  
    filled=True,  
    rounded=True,  
    fontsize=9  
)  
plt.show()
```



```
In [6]: from sklearn.model_selection import cross_val_score
```

```
cv_scores = cross_val_score(clf, X_encoded, y, cv=5, scoring="accuracy")
print("Cross-validation accuracies:", cv_scores)
print("Mean CV accuracy:", np.mean(cv_scores))
```

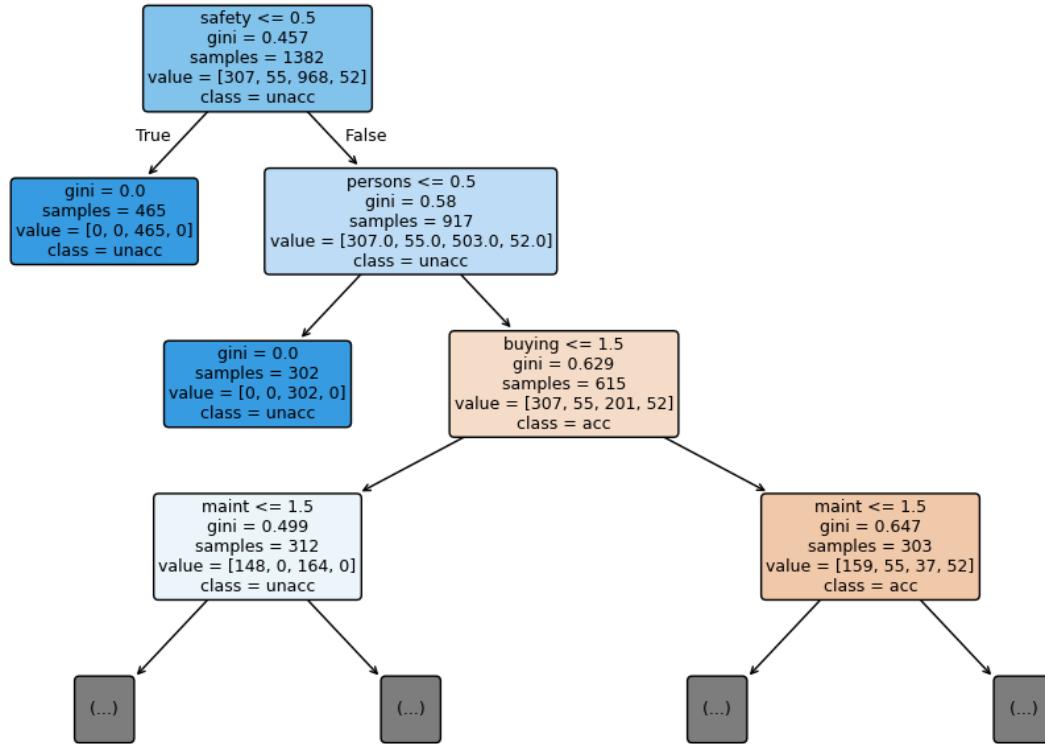
Cross-validation accuracies: [0.77745665 0.71098266 0.72543353 0.90434783 0.85217391]

Mean CV accuracy: 0.7940789143000754

```
In [7]: from sklearn import tree  
        import matplotlib.pyplot as plt
```

```
# Limit tree depth just for visualization (does not retrain model)
plt.figure(figsize=(12, 8))
tree.plot_tree(
    clf,
    feature_names=cols[:-1],
    class_names=clf.classes_,
    filled=True,
    rounded=True,
    fontsize=9,
    max_depth=3    # visualize only first 3 levels
)
plt.title("Decision Tree (First 3 Levels)")
plt.show()
```

Decision Tree (First 3 Levels)



Part-2: Churn Dataset

```
In [8]: import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split, StratifiedKFold, GridSearchCV
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder
from sklearn.pipeline import Pipeline
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt

RANDOM_STATE = 42
```

```
In [9]: path = "Telco-Customer-Churn_15081388-b1c7-4a07-a581-6b6a494e6346 (2).xlsx"
df = pd.read_excel(path)

print(df.shape)
df.head()
```

(7043, 21)

Out[9]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneServ
0	7590-VHVEG	Female	0	Yes	No	1	
1	5575-GNVDE	Male	0	No	No	34	
2	3668-QPYBK	Male	0	No	No	2	
3	7795-CFOCW	Male	0	No	No	45	
4	9237-HQITU	Female	0	No	No	2	

5 rows × 21 columns

In [10]:

```
# Try common target names
possible_targets = ["Churn", "churn", "CHURN"]
target_col = next((c for c in possible_targets if c in df.columns), None)
assert target_col is not None, "Couldn't find Churn column. Update target_col"

# Coerce typical numeric columns if present
for col in ["TotalCharges", "MonthlyCharges", "tenure"]:
    if col in df.columns:
        df[col] = pd.to_numeric(df[col], errors="coerce")

# Drop rows with missing target, and optionally impute features later
df = df.dropna(subset=[target_col]).copy()

# Trim whitespace in string columns
for c in df.select_dtypes(include="object").columns:
    df[c] = df[c].astype(str).str.strip()

# If 'customerID' exists and is an identifier, drop it
for id_col in ["customerID", "CustomerID", "customer_id"]:
    if id_col in df.columns:
        df = df.drop(columns=[id_col])

# Binary-size target if strings like 'Yes'/ 'No'
if df[target_col].dtype == 'object':
    if set(df[target_col].unique()) <= {"Yes", "No"}:
        df[target_col] = (df[target_col] == "Yes").astype(int)

y = df[target_col]
X = df.drop(columns=[target_col])

print("Target positive rate:", y.mean())
print("X shape:", X.shape)
```

Target positive rate: 0.2653698707936959

X shape: (7043, 19)

In [11]:

```
num_cols = X.select_dtypes(include=[np.number]).columns.tolist()
```

```

cat_cols = [c for c in X.columns if c not in num_cols]

# We'll treat 'Contract' as ordinal if present
ord_cols = []
ord_categories = []

if "Contract" in cat_cols:
    ord_cols.append("Contract")
    ord_categories.append(["Month-to-month", "One year", "Two year"])

# (Optional) You can add more ordinal mappings if you're confident about order
# if "tech_support_level" in cat_cols:
#     ord_cols.append("tech_support_level")
#     ord_categories.append(["None", "Basic", "Advanced"])

# Nominal columns are remaining categoricals minus ordinals
nom_cols = [c for c in cat_cols if c not in ord_cols]

print("Numeric:", num_cols)
print("Ordinal:", ord_cols)
print("Nominal:", nom_cols[:8], "... ({} nominal cols)".format(len(nom_cols)))

```

Numeric: ['SeniorCitizen', 'tenure', 'MonthlyCharges', 'TotalCharges']
Ordinal: ['Contract']
Nominal: ['gender', 'Partner', 'Dependents', 'PhoneService', 'MultipleLines',
'InternetService', 'OnlineSecurity', 'OnlineBackup'] ... (14 nominal cols)

```

In [13]: # Preprocessor (OrdinalEncoder for true ordinal, OneHot for nominal)
transformers = []

if ord_cols:
    transformers.append(("ord", OrdinalEncoder(categories=ord_categories, handle_unknown="ignore"), ord_cols))

if nom_cols:
    transformers.append(("onehot", OneHotEncoder(handle_unknown="ignore", sparse=False), nom_cols))

if num_cols:
    # passthrough numeric
    transformers.append(("num", "passthrough", num_cols))

preprocess = ColumnTransformer(transformers=transformers, remainder="drop", verbose=True)

transformers = []

if ord_cols:
    transformers.append(("ord", OrdinalEncoder(categories=ord_categories, handle_unknown="ignore"), ord_cols))

if nom_cols:
    transformers.append(("onehot", OneHotEncoder(handle_unknown="ignore", sparse=False), nom_cols))

if num_cols:
    # passthrough numeric
    transformers.append(("num", "passthrough", num_cols))

```

```
preprocess = ColumnTransformer(transformers=transformers, remainder="drop", vе  
In [14]: X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, stratify=y, random_state=RANDOM_STATE  
)
```

```
In [15]: full_tree = Pipeline([  
    ("prep", preprocess),  
    ("dt", DecisionTreeClassifier(  
        random_state=RANDOM_STATE, # fully grown: leave defaults (no max_depth)  
    ))  
])  
  
full_tree.fit(X_train, y_train)  
y_pred = full_tree.predict(X_test)  
  
print("Fully-grown tree - Test report")  
print(classification_report(y_test, y_pred, digits=4))
```

```
Fully-grown tree - Test report  
precision    recall   f1-score   support  
  
          0       0.8179     0.8029     0.8103      1035  
          1       0.4809     0.5053     0.4928       374  
  
accuracy                  0.7239      1409  
macro avg       0.6494     0.6541     0.6516      1409  
weighted avg     0.7285     0.7239     0.7261      1409
```

```
In [16]: # Pre-pruning via CV grid search  
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=RANDOM_STATE)  
  
param_grid = {  
    "dt__max_depth": [None, 5, 8, 12, 16, 20],  
    "dt__min_samples_split": [2, 10, 20],  
    "dt__min_samples_leaf": [1, 5, 10],  
    "dt__max_leaf_nodes": [None, 25, 50, 100],  
}  
  
prepruned = Pipeline([  
    ("prep", preprocess),  
    ("dt", DecisionTreeClassifier(random_state=RANDOM_STATE))  
])  
  
gs = GridSearchCV(  
    prepruned,  
    param_grid=param_grid,  
    scoring="f1" if len(np.unique(y))==2 else "accuracy",  
    cv=cv,  
    n_jobs=-1,  
    verbose=1  
)
```

```

gs.fit(X_train, y_train)

print("Best pre-pruning params:", gs.best_params_)
print("CV best score:", gs.best_score_)
print("Test score:", gs.best_estimator_.score(X_test, y_test))

y_pred_pre = gs.best_estimator_.predict(X_test)
print("\nPre-pruned tree - Test report")
print(classification_report(y_test, y_pred_pre, digits=4))

```

Fitting 5 folds for each of 216 candidates, totalling 1080 fits
 Best pre-pruning params: {'dt__max_depth': None, 'dt__max_leaf_nodes': 25, 'dt__min_samples_leaf': 10, 'dt__min_samples_split': 2}
 CV best score: 0.5841452146414278
 Test score: 0.7934705464868701

	precision	recall	f1-score	support
0	0.8509	0.8715	0.8611	1035
1	0.6189	0.5775	0.5975	374
accuracy			0.7935	1409
macro avg	0.7349	0.7245	0.7293	1409
weighted avg	0.7894	0.7935	0.7911	1409

In [17]: # Post-pruning: find α (cost-complexity pruning) with CV

```

# fitting an unpruned tree on TRAIN ONLY to get the  $\alpha$  path
tmp_model = Pipeline([
    ("prep", preprocess),
    ("dt", DecisionTreeClassifier(random_state=RANDOM_STATE))
])
tmp_model.fit(X_train, y_train)

# Get the underlying trained tree to compute pruning path:
# We need the transformed X_train to call cost_complexity_pruning_path
Xtr_trans = tmp_model.named_steps["prep"].transform(X_train)
ytr = y_train.values if hasattr(y_train, "values") else y_train

from sklearn.tree import DecisionTreeClassifier
base = DecisionTreeClassifier(random_state=RANDOM_STATE)
base.fit(Xtr_trans, ytr)
path = base.cost_complexity_pruning_path(Xtr_trans, ytr)
ccp_alphas = np.unique(path ccp_alphas) # increasing sequence

# Cross-validating over alpha values (wrap inside a fresh pipeline)
alphas_to_try = np.linspace(ccp_alphas.min(), ccp_alphas.max(), num=min(50, len(ccp_alphas)), endpoint=True)

post_scores = []
for a in alphas_to_try:
    model = Pipeline([
        ("prep", preprocess),
        ("dt", DecisionTreeClassifier(ccp_alpha=a, random_state=RANDOM_STATE))
    ])
    model.fit(Xtr_trans, ytr)
    post_scores.append(model.score(X_test, y_test))

```

```

        ("dt", DecisionTreeClassifier(random_state=RANDOM_STATE, ccp_alpha=a))
    ])
cv_score = cross_val_score(model, X_train, y_train, scoring="f1" if len(np
                           cv=cv, n_jobs=-1).mean()
post_scores.append(cv_score)

best_alpha = float(alphas_to_try[int(np.argmax(post_scores))])
best_alpha

```

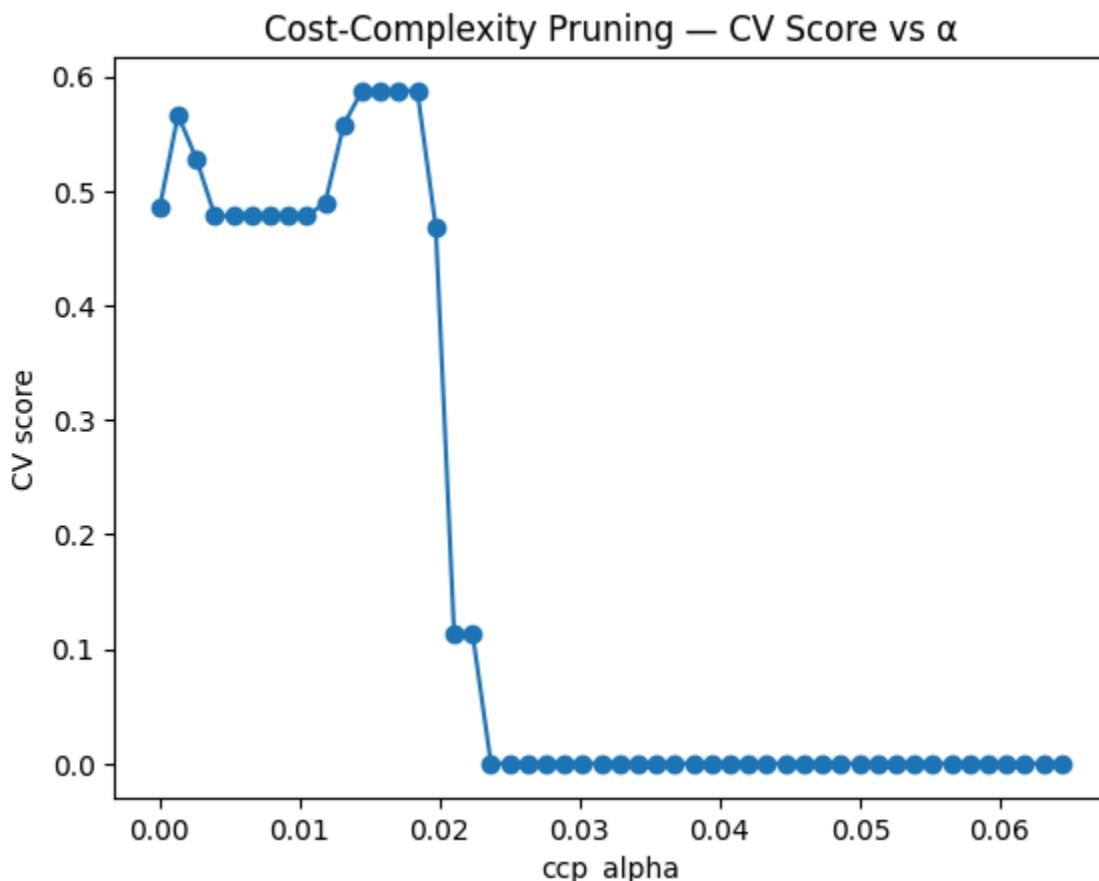
Out[17]: 0.014455730648082539

```

In [18]: # Plot CV score vs α (post-pruning curve)
plt.figure()
plt.plot(alphas_to_try, post_scores, marker="o")
plt.xlabel("ccp_alpha")
plt.ylabel("CV score")
plt.title("Cost-Complexity Pruning – CV Score vs α")
plt.show()

print("Selected α (ccp_alpha):", best_alpha)

```



Selected α (ccp_alpha): 0.014455730648082539

```

In [19]: # Train final post-pruned tree with best α and evaluate
postpruned = Pipeline([
    ("prep", preprocess),
    ("dt", DecisionTreeClassifier(random_state=RANDOM_STATE, ccp_alpha=best_alpha))
])

```

```
])
postpruned.fit(X_train, y_train)

y_pred_post = postpruned.predict(X_test)
print("Post-pruned tree – Test report")
print(classification_report(y_test, y_pred_post, digits=4))
```

```
Post-pruned tree – Test report
precision    recall   f1-score   support

          0       0.8462    0.8077    0.8265      1035
          1       0.5273    0.5936    0.5585      374

   accuracy                           0.7509      1409
macro avg       0.6867    0.7007    0.6925      1409
weighted avg    0.7615    0.7509    0.7554      1409
```

In []:

In []: